



A Inference: frame-by-frame gazing within each tile



B Downstream: merge positions, embed selected patches, then attend

Panel A outputs from all tiles



C Training only: measured reconstruction loss supervises the policy

